DEEP LEARNING ARCHITECTURES AND ATTENTION MECHANISMS

Section 1: Transformer Architecture Overview

The Transformer model, introduced in 2017 by Vaswani et al. in 'Attention Is All You Need', dispenses with recur

Section 2: Positional Encoding

Because the Transformer contains no recurrence and no convolution, in order for the model to make use of the o

Section 3: Retrieval-Augmented Generation (RAG)

Retrieval-Augmented Generation bridges pre-trained parametric memory with non-parametric external knowledg